

## C++ Strings

*Original handout written by Neal Kanodia, with Steve Jacobson.*

### C++ Strings

One of the most useful data types supplied in the C++ libraries is the `string`. A string is a variable that stores a sequence of letters or other characters, such as `"Hello"` or `"May 10th is my birthday!"`. Just like the other data types, to create a string we first declare it, then we can store a value in it. Strings in C++ are fairly similar to strings in Java, but there are also some important differences. Below, we provide a brief overview of using strings in C++.

We can declare strings and assign literal values to them as follows:

```
string testString;  
testString = "This is a string.";
```

We can combine these two statements into one line:

```
string testString = "This is a string.";
```

Often, we use strings as output, and `cout` works exactly like one would expect:

```
cout << testString << endl;
```

will print the same result as

```
cout << "This is a string." << endl;
```

In order to use the string data type, the C++ string header file must be included at the top of the program. Also, you'll need to include `genlib.h` to make the short name `string` visible instead of requiring the cumbersome `std::string`. (As a side note, `std` is a C++ *namespace* for many pieces of functionality that are provided in standard C++ libraries. For the purposes of this class, you won't need to know about namespaces.) Thus, you would have the following `#include`'s in your program in order to use the `string` type.

```
#include <string>  
#include "genlib.h"
```

#### **Note for students who programmed in C previously, but not C++**

Throughout this handout, there will be boxed sections that will help you transition from C-style `char*` strings to C++ strings.

The first and foremost thing to keep in mind is that C++ strings are a completely new data type. C++ strings are not pointers to character arrays, you can not use direct pointer-style manipulations on them, and most techniques you used to manipulate C-style strings are not applicable. However, comparable (and in many cases, better) facilities are supported.

Memory management is one of the largest differences between C and C++ strings. In C, the programmer is expected to allocate and deallocate memory for strings, and avoid writing outside the allocated space. In C++, all of this will be taken care of automatically! For example, in C++, strings are automatically copied for string assignment or when passed as parameters. In many ways, a C++ `string` is as simple to use as an `int`, `double`, or `char`.

## Basic Operations

Let's go into specifics about the string manipulations you'll be doing the most.

**Counting the number of characters in a string.** The `.length()` function returns the number of characters in a string, including spaces and punctuation. Like many of the string operations, `length` is a *member function*, and we invoke member functions using *dot notation*. The string that is the receiver is to the left of the dot, the member function we are invoking is to the right, (e.g. `str.length()`). In such an expression, we are requesting the length from the variable `str`.

*example program:*

```
#include <string>
#include <iostream>
#include "genlib.h"

int main() {
    string small, large;
    small = "I am short";
    large = "I, friend, am a long and elaborate string indeed";

    cout << "The short string is " << small.length()
          << " characters." << endl;
    cout << "The long string is " << large.length()
          << " characters." << endl;
    return 0;
}
```

*output:*

```
The short string is 10 characters.
The long string is 48 characters.
```

**Accessing individual characters.** Using square brackets, you can access individual characters within a string, similar to Java. Just as with array access, positions are numbered starting from 0, and the last position is (length - 1). Using square brackets, you can both read the character at a position and assign to that position. Note that the ability to assign characters directly into a string in C++ makes them different than in Java.

*example program:*

```
#include <string>
#include <iostream>
#include "genlib.h"

int main() {
    string test;
    test = "I am Q the omnipot3nt";

    char ch = test[5];          // ch is 'Q'
    test[18] = 'e';             // we correct misspelling of omnipotent

    cout << test << endl;
    cout << "ch = " << ch << endl;
    return 0;
}
```

*output:*

```
I am Q the omnipotent
ch = Q
```

Be careful not to access positions outside the bounds of the string! The square bracket operator is not range-checked (for efficiency reasons) and thus reading from or writing to an out-of-bounds index tends to produce difficult-to-track-down errors. There is an alternate member function `.at(index)` that retrieves the character at a position with the benefit of built-in range-checking.

**Note for former C programmers**

C-style strings use a `'\0'` character to mark the end of a string. C++ strings do not have such an end marker. The last character of C++ string is at position `length()-1`.

**Passing, returning, assigning strings.** C++ string variables are designed to behave like ordinary primitive types with regard to assignment. Assigning one string to another makes a new copy of the character sequence. For example:

```
string str1= "hello";
string str2 = str1;      // makes a new copy
str1[0] = 'y';          // changes only str1, str2 is not changed
```

Similarly, passing and returning strings from functions copies the string contents. If you change a string parameter within a function, changes are not seen in the calling function unless you have specifically passed the string by reference.

**Note for former C programmers**

This is completely different than C-style strings. For `char*` strings, string assignment, passing, and return only copy the *pointer*, thus aliasing the contents. This can lead to unexpected consequences. C++ strings remove this weirdness, each C++ string is completely independent of all others. When C++ strings are assigned, passed into or returned from functions, a new copy of the string is made.

**Comparing two strings:** You can compare two strings for equality using the ordinary `==` and `!=` operators (just like they were `ints`). Suppose you ask the user for his or her name. If the user is Julie, the program prints a warm welcome. If the user is not Neal, the program prints the normal message. Finally... if the user is Neal, it prints a less enthusiastic response.

*example program:*

```

#include <string>
#include <iostream>
#include "simpio.h"
#include "genlib.h"

int main() {
    string myName = "Neal";

    while (true) {
        cout << "Enter your name (or 'quit' to exit): ";
        string userName = GetLine();

        if (userName == "Julie") {
            cout << "Hi, Julie! Welcome back!" << endl << endl;
        } else if (userName == "quit") {
            // user is sick of entering names, so let's quit
            cout << endl;
            break;
        } else if (userName != myName) {
            // user did not enter quit, Julie, or Neal
            cout << "Hello, " << userName << endl << endl;
        } else {
            cout << "Oh, it's you, " << myName << endl << endl;
        }
    }

    return 0;
}

```

*output:*

```

Enter your name (or 'quit' to exit): Neal
Oh, it's you, Neal

Enter your name (or 'quit' to exit): Julie
Hi, Julie! Welcome back!

Enter your name (or 'quit' to exit): Leland
Hello, Leland

Enter your name (or 'quit' to exit): quit

Press any key to continue

```

You can use `<`, `<=`, `>`, and `>=` to order strings. These operators compare strings lexicographically, character by character and are case-sensitive. The following comparisons all evaluate to true: `"A" < "B"`, `"App" < "Apple"`, `"help" > "hello"`, `"Apple" < "apple"`. The last one might be a bit confusing, but the ASCII value for 'A' is 65, and comes before 'a', whose ASCII value is 97. So "Apple" comes before "apple" (or, for that matter, any other word that starts with a lower-case letter).

**Note for former C programmers**

You know that `==` on `char*` (C-style) strings just compares the two pointers for equality (not very useful!). But `==` on C++ strings is much more useful since it truly compares the characters in the strings for equality.

**Appending to a string:** C++ strings are wondrous things. Suppose you have two strings, `s1` and `s2` and you want to create a new string of their concatenation. Conveniently, you can just write `s1 + s2`, and you'll get the result you'd expect. Similarly, if you want to append to the end of string, you can use the `+=` operator. You can append either another string or a single character to the end of a string.

**example program:**

```
#include <string>
#include <iostream>
#include "genlib.h"

int main() {
    string firstname = "Leland";
    string lastname = " Stanford";

    string fullname = firstname + lastname; // concat the two strings

    fullname += ", Jr";           // append another string
    fullname += '.';              // append a single char

    cout << firstname << lastname << endl;
    cout << fullname << endl;

    return 0;
}
```

**output:**

```
Leland Stanford
Leland Stanford, Jr.
```

**More (Less Used) Operations**

The string class has many more operations; we'll show just a few of the more useful ones below.

**Searching within a string.** The string member function `find` is used to search within a string for a particular string or character. A sample usage such as `str.find(key)` searches the receiver string `str` for the `key`. The parameter `key` can either be a string or a character. (We say the `find` member function is *overloaded* to allow more than one usage). The return value is either the starting *position* where the key was found or the constant `string::npos` which indicates the key was not found.

Occasionally, you'll want to control what part of the string is searched, such as to find a second occurrence past the first. There is an optional second integer argument to `find` which allows you to specify the starting position; when this argument is not given, 0 is assumed. Thus,

`str.find(key, n)` starts at position `n` within `str` and will attempt to find `key` from that point on. The following code should make this slightly clearer:

*example program:*

```
#include <string>
#include <iostream>
#include "genlib.h"

int main() {
    string sentence = "Yes, we went to Gates after we left the dorm.";

    int firstWe = sentence.find("we");           // finds the first "we"
    int secondWe = sentence.find("we", firstWe+1); // finds "we" in "went"
    int thirdWe = sentence.find("we", secondWe+1); // finds the last "we"
    int gPos = sentence.find('G');
    int zPos = sentence.find('Z');               // returns string::npos

    cout << "First we: " << firstWe << endl;
    cout << "Second we: " << secondWe << endl;
    cout << "Third we: " << thirdWe << endl;

    cout << "Is G there? ";
    if (gPos == string::npos) {
        cout << "No!" << endl;
    } else {
        cout << "Yes!" << endl;
    }
    cout << "Is Z there? ";
    if (zPos == string::npos) {
        cout << "No!" << endl;
    } else {
        cout << "Yes!" << endl;
    }
    return 0;
}
```

*output:*

```
First we: 5
Second we: 8
Third we: 28
Is G there? Yes!
Is Z there? No!
```

**Extracting substrings.** Sometimes you would like to create new strings by excerpting portions of a string. The `substr` member function creates such substrings from pieces of the receiver string. You specify the starting position and the number of characters. For example, `str.substr(start, length)` returns a new string consisting of the characters from `str` starting at the position `start` and continuing for `length` characters. Invoking this member function does not change the receiver string, it makes a *new* string with a copy of the characters specified.

*example program:*

```

#include <string>
#include <iostream>
#include "genlib.h"

int main() {
    string oldSentence;
    oldSentence = "The quick brown fox jumped WAY over the lazy dog";
    int len = oldSentence.length();

    cout << "Original sentence: " << oldSentence << endl;

    // Create a new sentence without "WAY ":

    // First, find the position where WAY occurs
    int i = oldSentence.find("WAY ");

    // Second, get the characters up to "WAY "
    string newSentence = oldSentence.substr(0, i);

    cout << "Modified sentence: " << newSentence << endl;

    // Finally, append the characters after "WAY "
    //
    // We use a special property of .substr here to make our life
    // easier. If we don't specify the length, it just gets the
    // rest of the string.
    newSentence += oldSentence.substr(i + 4);

    cout << "Completed sentence: " << newSentence << endl;

    return 0;
}

```

*output:*

```

Original sentence: The quick brown fox jumped WAY over the lazy dog
Modified sentence: The quick brown fox jumped
Completed sentence: The quick brown fox jumped over the lazy dog

```

There are a couple of special cases for `.substr(start, length)`. If `start` is negative, it will cause a run-time error. If `start` is past the end of the string it will return an empty string (""). If `length` is longer than the number of characters from the start position to the end of the string, it truncates to the end of the string. If `length` is negative, then the behavior is undefined, so make sure that `length` is always non-negative. If you leave off the second argument, the number of characters from the starting position to the end of the receiver string is assumed.

**Modifying a string by inserting and replacing:** Finally, let's cover two other useful member functions that modify the receiver string. The first, `str1.insert(start, str2)`, inserts `str2` at position `start` within `str1`, shifting the remaining characters of `str1` over. The second, `str1.replace(start, length, str2)`, removes from `str1` a total of `length` characters starting at the position `start`, replacing them with a copy of `str2`. It is important to note that these member functions do modify the receiver string.

*example program:*

```

#include <string>
#include <iostream>
#include "genlib.h"

int main() {
    int start_pos;
    string sentence = "CS 106B sucks.";
    cout << sentence << endl;

    // Insert "kind of" at position 8 in sentence
    sentence.insert(8, "kind of ");
    cout << sentence << endl;

    // Replace the 10 characters "kind of su"
    // with the string "ro" in sentence
    sentence.replace(8, 10, "ro");
    cout << sentence << endl;

    return 0;
}

```

*output:*

```

CS 106B sucks.
CS 106B kind of sucks.
CS 106B rocks.

```

**Obtaining a C-style `char*` from a string**

Remember, a C++ string is not the same thing as a C-style string (which is merely a `char*` pointer to a sequence of characters terminated by a null character `\0`). Although old-style C `char*` strings and C++ strings can co-exist in a program, almost all our use will be of C++ strings, since they have a much richer set of operations and are less error-prone to work with. I say "almost always" because in a few unavoidable situations, we are forced to use old-style C strings, most notably when working with file streams. We can convert a C++ string to the old-style representation using the `.c_str()` member function. One use we will see of this is to get a `char*` to pass to the `iostream` `open` function.

*example program:*

```

#include <string>
#include <iostream>
#include <fstream>
#include "genlib.h"

int main() {
    ifstream fs;
    string s, filename = "courseinfo.txt";

    // open member function requires a C-style string, must convert!
    fs.open(filename.c_str());

    if (fs.fail()) return -1; // could not open the file!
    // rest of program here...
    fs.close();
    return 0;
}

```



## The CS106 Library: `strutils.h`

In addition to the standard library support for strings, there are a few extensions that the CS106 libraries provide. To use these functions, the `strutils.h` library must be included:

```
#include <strutils.h>
```

**IntegerToString, RealToString, StringToInteger, StringToReal:** Often your programs will need to convert a string to a number or vice versa. These functions do just that, with the 'integer' functions operating on `int` and the 'real' functions on `double`.

### *example program:*

```
#include <string>
#include <iostream>
#include "strutils.h"
#include "genlib.h"

int main() {
    string str1 = "5.6";

    double num = StringToReal(str1);
    string str2 = IntegerToString(45);

    cout << "The original string is " << str1 << "." << endl;
    cout << "The number is " << num << "." << endl;
    cout << "The new string is " << str2 << "." << endl;

    return 0;
}
```

### *output:*

```
The original string is 5.6.
The number is 5.6
The new string is 45.
```

Any integer or real can be safely converted to a string. However, converting in the other direction, if you pass an improperly formatted string to convert to an integer or real, an error is raised by the conversion functions.

**ConvertToLowerCase, ConvertToUpperCase:** These functions take a string, and return a new string with all letters in lower or upper case respectively. These can be used to change two strings to a uniform case before comparing to allow a case-insensitive comparison.

*example program:*

```
#include <string>
#include <iostream>
#include "strutils.h"
#include "genlib.h"

int main()
{
    string appleFruit = "apples";
    string orangeFruit = "ORANGES";

    cout << "Do " << appleFruit << " come before "
          << orangeFruit << "? ";
    if (appleFruit < orangeFruit)
        cout << "Yes!" << endl;
    else
        cout << "Nope...." << endl;

    string lowerOrangeFruit = ConvertToLowerCase(orangeFruit);

    cout << "Do " << appleFruit << " come before "
          << lowerOrangeFruit << "? ";
    if (appleFruit < lowerOrangeFruit)
        cout << "Yes!" << endl;
    else
        cout << "Nope...." << endl;

    return 0;
}
```

*output:*

```
Do apples come before ORANGES?  Nope....
Do apples come before oranges?  Yes!
```